

Architectural Taxonomy & Latency Spectrum for Spoken Conversational Intelligence

1. Cascaded Discrete Pipeline (ASR + LLM + TTS)



2. End-to-End Duplex Audio LLMs (Moshi, Mini-Omni, Llama-Omni)



3. GuimLab Continuous Neuromorphic Substrate (Pure CUDA/C++20)

